Operational Intelligence Publication Pack

Shareable navigation layer for the doctrine, reference architecture, Operations Room, diagrams, tables, and examples.

Navigation

Asset	Purpose
Canonical Doctrine	Defines why Operational Intelligence exists, boundaries, glossary, and core thesis.
Reference Architecture	Defines implementation-neutral contracts, schemas, state machines, evaluation gates, and conformance.
Operations Room	Interactive synthetic artifact demonstrating the doctrine through OI-ROOM-001.
Diagram Pack	Architecture, state-machine, sequence, evidence graph, and replay-loop diagrams.
Comparison Tables	Adjacent discipline comparison and claim classification.
Decision Packet	Concrete review packet for an evidence-backed recommendation.
Printable Walkthrough	Inspect OI-ROOM-001 as evidence, hypotheses, gates, and operator decision.
Glossary Card	Canonical terms for writers, reviewers, and retrieval.
Evidence Pack	Benchmark rubric, control comparisons, practitioner review, evidence ledger, and falsification criteria.

Architecture Flow

Signal -> Transaction -> Topology -> Evidence -> Reasoning -> Memory -> Evaluation -> Decision -> Operator -> Learning. Learning feeds reviewed memory and future evaluation fixtures.

Ten-layer operational reasoning path

Signal ▶ **Transaction** ▶ **Topology** ▶ **Evidence** ▶ **Reasoning**

Memory ▶ **Evaluation** ▶ **Decision** ▶ **Operator** ▶ **Learning**

Evaluation Rule

No aggregate trust score replaces dimension-level results. Retrieval, grounding, citation, refusal, evidence attribution, hypothesis quality, contradiction handling, replay quality, latency, operator approval, and recommendation safety must be evaluated separately.

Evidence Question

The question that moves the model forward is: what evidence would convince another experienced engineer that this model is useful? The answer should come from implementation examples, benchmarks, case studies, evaluations, practitioner feedback, and backward-compatible refinements.

Evidence Graph Shape

A publication-ready Operational Intelligence artifact must show how evidence supports, contradicts, limits, or informs a hypothesis. The graph below uses only the synthetic OI-ROOM-001 case.

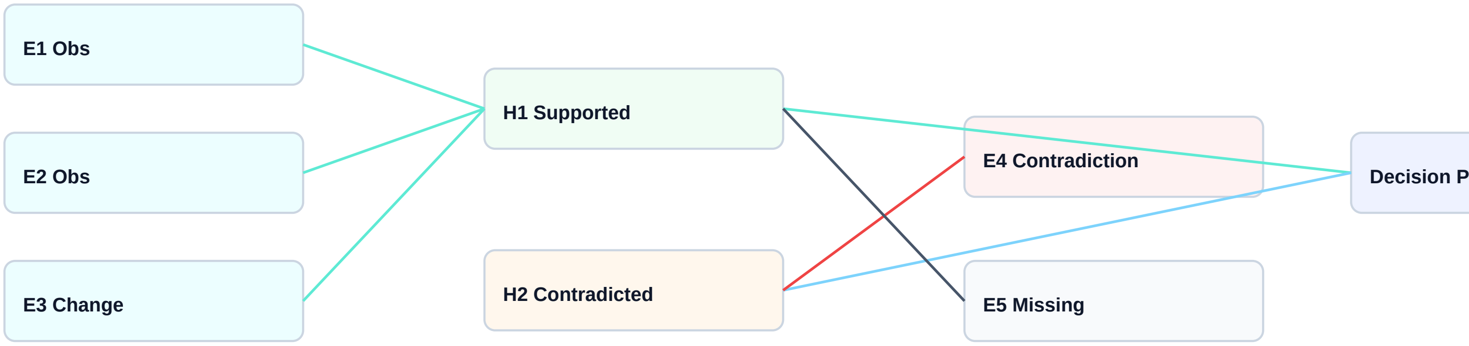

Observation, contradiction, and missing evidence remain visible before a reviewable decision packet is drafted.

Decision Handoff

Decision packet handoff: the assistant prepares a bounded packet; accountable humans approve, reject, escalate, or request evidence.

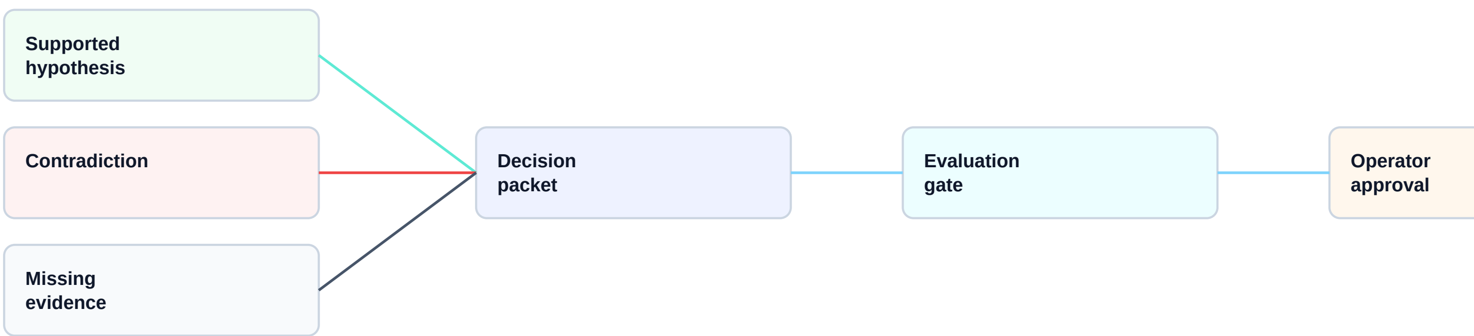

Decision Packet Structure

Decision Packet Field	Required content
Recommendation	Action framed as a reviewable option, not autonomous execution.
Evidence	Supporting, contradictory, and missing evidence cited separately.
Risk	Blast radius, reversibility, uncertainty, and owner named.
Approval	Operator, approval class, fallback, and escalation path.
Learning	Replay seed, eval fixture, and memory update after review.

Cross-reference rule

Definitions and boundaries should point to the Canonical Doctrine. Implementation behavior should point to the Reference Architecture. Shareable diagrams, tables, walkthroughs, and PDFs should point to the Publication Pack. Proof, benchmarks, practitioner review, and falsification should point to the Evidence Pack.